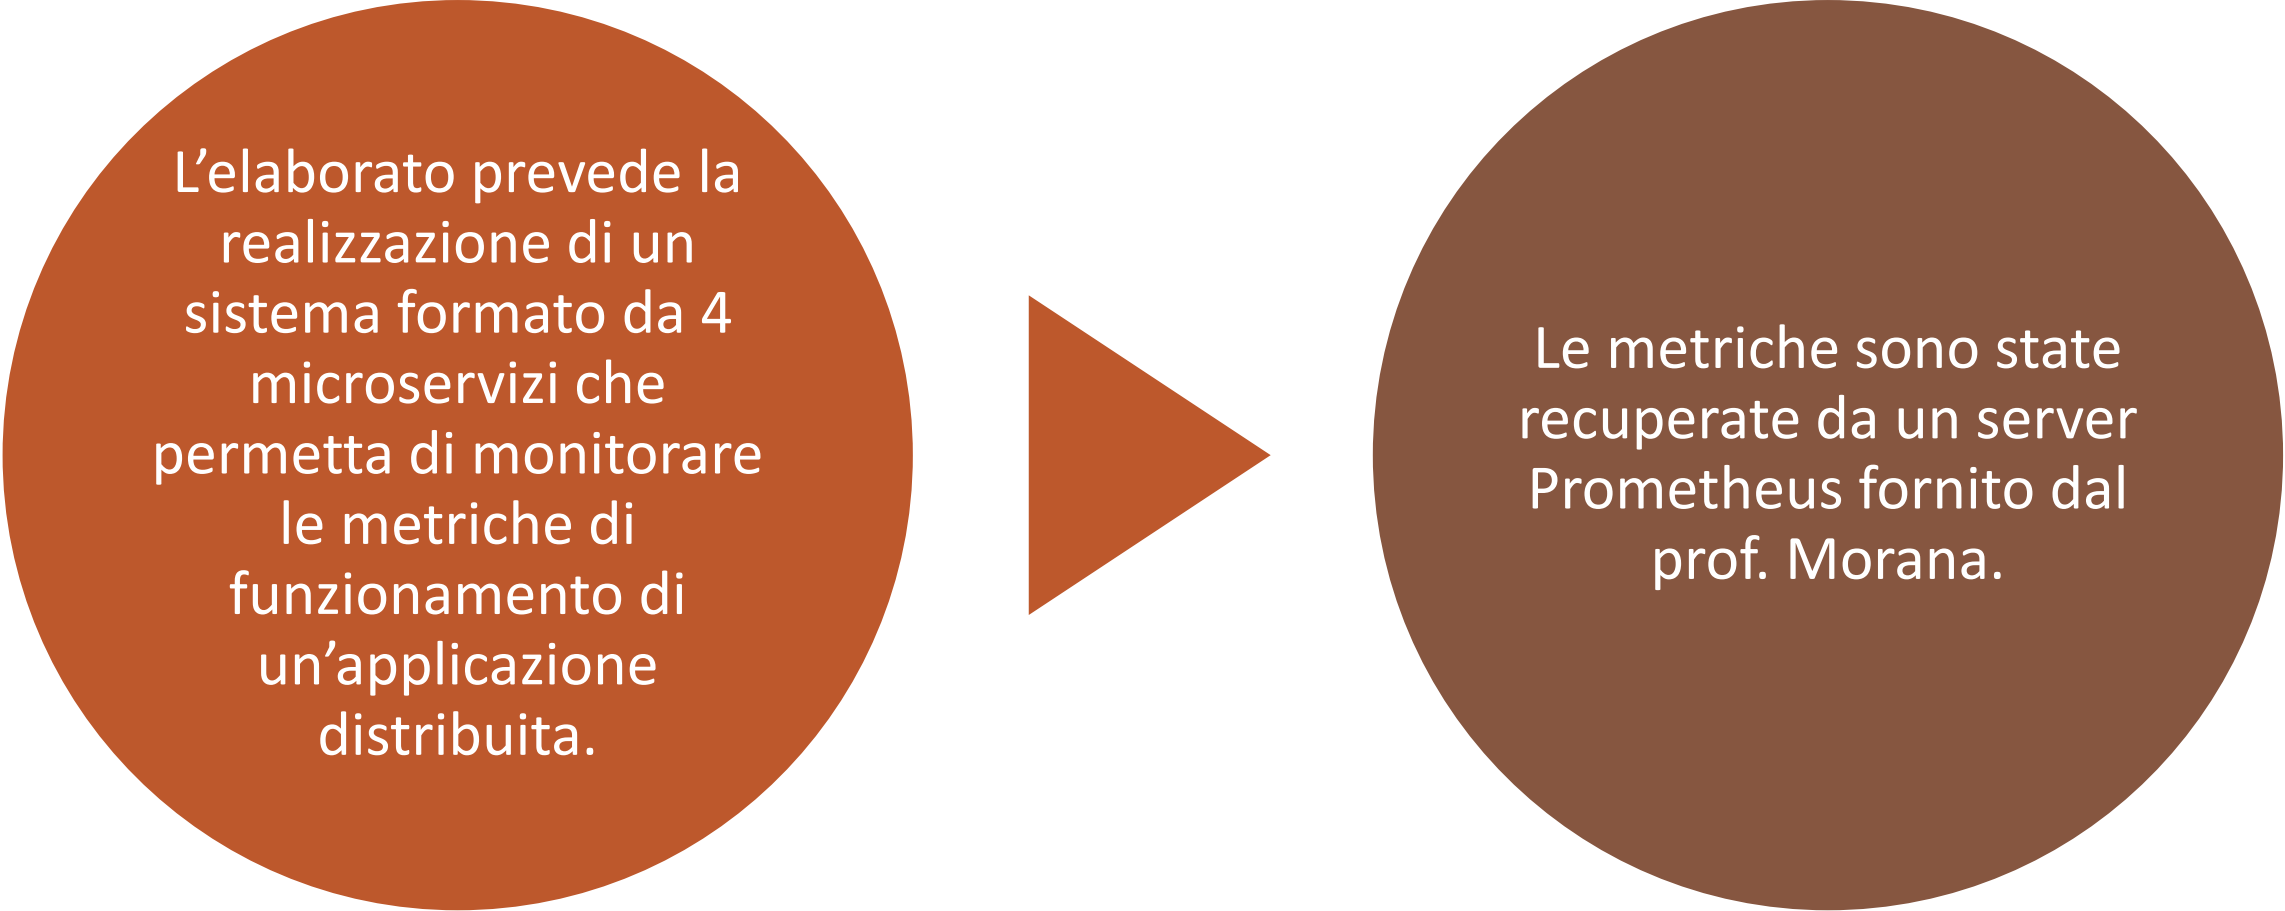


Corso di Distributed Systems & Big Data

Presentazione Progetto in itinere

SVOLTO DA:
LEONE DAMIANO
SPADARO SAPARI LORENZO

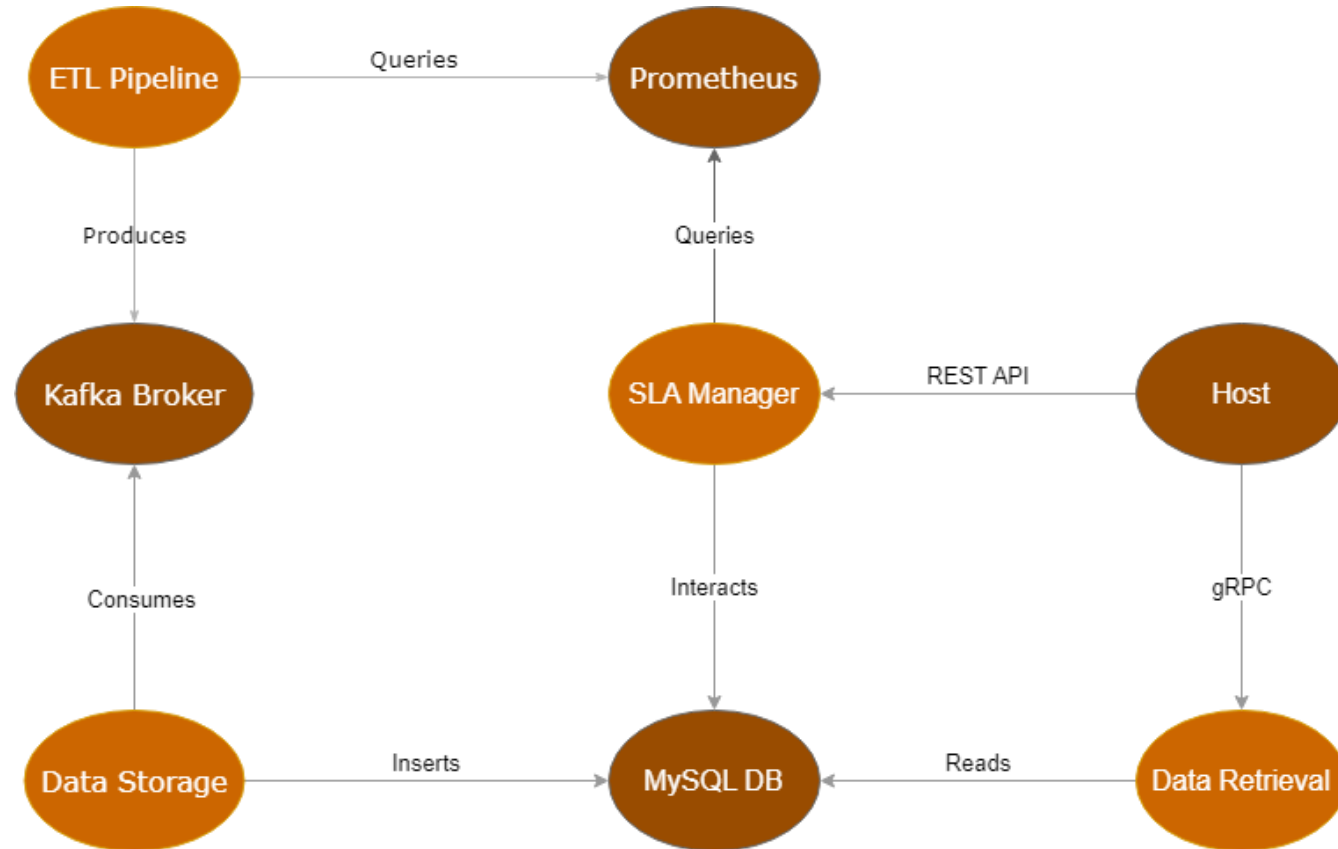
Contenuto elaborato



L'elaborato prevede la realizzazione di un sistema formato da 4 microservizi che permetta di monitorare le metriche di funzionamento di un'applicazione distribuita.

Le metriche sono state recuperate da un server Prometheus fornito dal prof. Morana.

Diagramma del Sistema



Metriche selezionate

Tra le molteplici metriche esposte dal server Prometheus fornito, le 5 metriche scelte per la formazione di un Service Level Agreement sono:

- `node_memory_memFree_bytes`: Byte di Memoria Liberi (non include cache e buffer)
- `node_memory_memAvailable_bytes`: Memoria disponibile (include memoria cache e buffers)
- `node_filesystem_free_bytes`: Spazio libero su disco (include lo spazio riservato per root)
- `node_filesystem_avail_bytes`: Spazio libero su disco (spazio disponibile come utente user)
- `node_filefd_allocated`: Numero totale filedescriptors allocati

Queste metriche fanno parte del job 'host' (Node Exporter).

Data Pipeline

```
def elaboraMetrica(nome_metrica,metrica,partizione,producer):
    isStationary=False
    critValues=[]
    metrics_data={}
    cyclical_component={}
    seasonal_component={}
    trend_component={}

    #Augmented dickey fuller test per valutare la staticità
    result_adf = adfuller(metrica,maxlag=3)
    out = pd.Series(result_adf[0:4],index=['ADF test statistic','p-value','# lags used','# observations'])
    if result_adf[1] <0.05:
        isStationary=True

    for key,val in result_adf[4].items():
        critValues.append(val)

    #filtro hodrick-prescott per prendere informazioni sulla ciclicità
    cyclical, trend = hpfilter(metrica, lamb=1600)
    cyclical.keys=cyclical.keys().strftime('%Y-%m-%d %H:%M:%S')

    #Autocorrelazione
    result_acf= acf(metrica)
    np.nan_to_num(result_acf,copy=False)
```

```
#valutazione metriche nelle ultime 1,3,12h
hourly_data = metrica.resample('1H').last().agg(['max', 'min', 'mean','std']).to_dict()
three_hourly_data = metrica.resample('3H').last().agg(['max', 'min', 'mean','std']).to_dict()
twelve_hourly_data = metrica.resample('12H').last().agg(['max', 'min', 'mean','std']).to_dict()

tsr = trend.resample(rule='1T').mean()

#predizione delle metriche per i prox 10 min
tsmodel = ExponentialSmoothing(tsr, trend='mul', seasonal='add',seasonal_periods=550).fit()
prediction = tsmodel.forecast(10)
predicted_metrics=prediction.agg(['max','min','mean']).to_dict()

#invio dati al broker kafka
total_data={'metric_name':nome_metrica,partizione: {'values':metrics_data,'adfuller_statistic':result_adf[0],'adfuller_p_val':result_adf[1],
            'adfuller_stationary':isStationary,'adfuller_critical':critValues,'acf':result_acf.tolist(),
            'decompose_season':seasonal_component,'decompose_trend':trend_component,'cyclical_component':cyclical_component,
            'hourly_data':hourly_data,'three_hourly_data':three_hourly_data,'twelve_hourly_data':twelve_hourly_data,
            '10m_prediction':predicted_metrics}}

producer.send("prometheusdata",total_data)
producer.flush()
```

Data Storage

```
cnx = mysql.connector.connect(host = "db",database= "test_dsbd",user = "user",password = "password")
cursor = cnx.cursor()

consumer=KafkaConsumer(bootstrap_servers='kafka:9092',api_version=(0,10,2),auto_offset_reset='earliest',enable_auto_commit=False)
consumer.subscribe(['prometheusdata'])

# Read data from kafka
for message in consumer:
    data=json.loads(message[6])
    metric_name=data['metric_name']
    for disk, values in data.items():
        if disk == 'metric_name':
            continue

        sampled_values=values['values']

        adfuller_statistic = values['adfuller_statistic']
        adfuller_p_value = values['adfuller_p_val']
        adfuller_stationary = values['adfuller_stationary']
        adfuller_critical= values['adfuller_critical']

        acf= values['acf']
```

```
hourly_data = values['hourly_data']
hourly_max=str(hourly_data['value']['max'])
hourly_min= str(hourly_data['value']['min'])
hourly_mean=str(hourly_data['value']['mean'])
hourly_std= str(hourly_data['value']['std'])

three_hourly_data = values['three_hourly_data']
three_hourly_max= str(three_hourly_data['value']['max'])
three_hourly_min= str(three_hourly_data['value']['min'])
three_hourly_mean=str(three_hourly_data['value']['mean'])
three_hourly_std= str(three_hourly_data['value']['std'])

twelve_hourly_data = values['twelve_hourly_data']
twelve_hourly_max= str(twelve_hourly_data['value']['max'])
twelve_hourly_min= str(twelve_hourly_data['value']['min'])
twelve_hourly_mean=str(twelve_hourly_data['value']['mean'])
twelve_hourly_std= str(twelve_hourly_data['value']['std'])

prediction = values['10m_prediction']
prediction_max= str(prediction['max'])
prediction_min= str(prediction['min'])
```

Data Storage cont.

```
query6="INSERT INTO aggregates(nome,partizione,intervallo,tipologia,value) VALUES(%s,%s,%s,%s,%s)"
tuples=[]
tuples.append([metric_name,disk,'1h','max',hourly_max])
tuples.append([metric_name,disk,'1h','min',hourly_min])
tuples.append([metric_name,disk,'1h','mean',hourly_mean])
tuples.append([metric_name,disk,'1h','std_dev',hourly_std])
tuples.append([metric_name,disk,'3h','max', three_hourly_max])
tuples.append([metric_name,disk,'3h','min',three_hourly_min])
tuples.append([metric_name,disk,'3h','mean',three_hourly_mean])
tuples.append([metric_name,disk,'3h','std_dev',three_hourly_std])
tuples.append([metric_name,disk,'12h','max',twelve_hourly_max])
tuples.append([metric_name,disk,'12h','min',twelve_hourly_min])
tuples.append([metric_name,disk,'12h','mean',twelve_hourly_mean])
tuples.append([metric_name,disk,'12h','std_dev',twelve_hourly_std])

for tup in tuples:
    cursor.execute(query6, tup)

query7="INSERT INTO aggregates(nome,partizione,intervallo,tipologia,valore_predetto) VALUES(%s,%s,%s,%s,%s)"
value7=(metric_name,disk,'10m','max',prediction_max)
cursor.execute(query7, value7)
value8=(metric_name,disk,'10m','min',prediction_min)
cursor.execute(query7, value8)
value9=(metric_name,disk,'10m','mean',prediction_mean)
cursor.execute(query7, value9)
cnx.commit()
```

Data Retrieval — .proto file

```
1 syntax = "proto3";
2
3 package dataretrieval;
4
5
6 //Run in the working directory: python -m grpc_tools.protoc -I ./ --python_out=. --pyi_out=. --grpc_python_out=. ./retrieval.proto
7
8 service RetrievalService {
9
10     //unary rpc
11     rpc GetDickeyFuller(RetrievalRequest) returns (RetrievalDickeyFuller) {}
12
13     //client-unary, server-streaming rpc
14     rpc ListDecomposition(RetrievalRequest) returns (stream RetrievalDecomposition) {}
15
16     rpc ListHodrickPrescott(RetrievalRequest) returns (stream RetrievalHodrickPrescott) {}
17
18     rpc ListAutocorrelation(RetrievalRequest) returns (stream RetrievalAutocorrelation) {}
19
20     rpc ListAggregates(RetrievalRequest) returns (stream RetrievalAggregates) {}
21 }
22
23 // Il messaggio di richiesta conterrà la query.
24 message RetrievalRequest {
25     string query = 1;
26 }
27
28 // Il messaggio di risposta conterrà i dati in modo strutturato in base alla tabella interrogata.
29 message RetrievalDickeyFuller {
30     string name = 1;
31     string partition = 2;
32     float test_statistic = 3;
33     float p_value = 4;
34     bool is_stationary = 5;
35     float crit_value1 = 6;
36     float crit_value5 = 7;
37     float crit_value10 = 8;
38 }
```

```
40 message RetrievalDecomposition {
41     string name = 1;
42     string partition = 2;
43     string timestamp = 3;
44     string typology = 4;
45     float value = 5;
46 }
47
48 message RetrievalHodrickPrescott {
49     string name = 1;
50     string partition = 2;
51     string timestamp = 3;
52     float value = 4;
53 }
54
55 message RetrievalAutocorrelation {
56     string name = 1;
57     string partition = 2;
58     float value = 3;
59 }
60
61 message RetrievalAggregates {
62     string name = 1;
63     string partition = 2;
64     string range = 3;
65     string typology = 4;
66     float value = 5;
67     float predicted_value = 6;
68 }
```


Data Retrieval - asynch_client.py

```
250 async def run() -> None:
251     async with grpc.aio.insecure_channel('server_data_retrieval:50051') as channel:
252         stub = retrieval_pb2_grpc.RetrievalServiceStub(channel)
253
254     while(True):
255         user_input = int(input("Inserisci un numero intero corrispondente a una metrica:\n\n1-node_memory_MemFree_bytes\n2-node_memory_memAvail_bytes\n3-node_filesystem_free_bytes\n4-node_filesystem_avail_bytes\n5-node_filefd_allocated\n6-Richiedi i dati relativi a tutte le metriche\n7-Exit\n\n"))
256
257         match user_input:
258
259             case 1:
260                 await node_memory_MemFree_bytes(stub)
261
262             case 2:
263                 await node_memory_MemAvailable_bytes(stub)
264
265             case 3:
266                 await node_filesystem_free_bytes_dev_sda2(stub)
267                 await node_filesystem_free_bytes_tmpfs(stub)
268
269             case 4:
270                 await node_filesystem_avail_bytes_dev_sda2(stub)
271                 await node_filesystem_avail_bytes_tmpfs(stub)
272
273             case 5:
274                 await node_filefd_allocated(stub)
275
276             case 6:
277                 await node_memory_MemFree_bytes(stub)
278                 await node_memory_MemAvailable_bytes(stub)
279                 await node_filesystem_free_bytes_dev_sda2(stub)
280                 await node_filesystem_free_bytes_tmpfs(stub)
281                 await node_filesystem_avail_bytes_dev_sda2(stub)
282                 await node_filesystem_avail_bytes_tmpfs(stub)
283                 await node_filefd_allocated(stub)
284
285             case 7:
286                 sys.exit(0)
287
288             case _:
289                 print("Errore! Inserire un numero intero compreso tra 1 e 7.\n")
290
```

Inserisci un numero intero corrispondente a una metrica:

```
1-node_memory_MemFree_bytes
2-node_memory_memAvail_bytes
3-node_filesystem_free_bytes
4-node_filesystem_avail_bytes
5-node_filefd_allocated
6-Richiedi i dati relativi a tutte le metriche
7-Exit
>>
```

async_client.py cont. - esempio metrica 1

```
67 async def node_memory_MemFree_bytes(stub: retrieval_pb2_grpc.RetrievalServiceStub) -> None:
68     print( "Metrica 1: node_memory_MemFree_bytes\n")
69
70     request = retrieval_pb2.RetrievalRequest(query='select * from dickey_fuller where nome="memory_MemFree_bytes"')
71     print("----- GetDickeyFuller ----- \n")
72     await get_dickey_fuller(stub, request)
73
74     request = retrieval_pb2.RetrievalRequest(query='select * from decomposition where nome="memory_MemFree_bytes"')
75     print("----- ListDecomposition ----- \n")
76     await list_decomposition(stub, request)
77
78     request = retrieval_pb2.RetrievalRequest(query='select * from hodrick_prescott where nome="memory_MemFree_bytes"')
79     print("----- ListHodrickPrescott ----- \n")
80     await list_hodrick_prescott(stub, request)
81
82     request = retrieval_pb2.RetrievalRequest(query='select * from autocorrelation where nome="memory_MemFree_bytes"')
83     print("----- ListAutocorrelation ----- \n")
84     await list_autocorrelation(stub, request)
85
86     request = retrieval_pb2.RetrievalRequest(query='select * from aggregates where nome="memory_MemFree_bytes"')
87     print("----- ListAggregates ----- \n")
88     await list_aggregates(stub, request)
89
90     print( "Fine stampa Metrica 1: node_memory_MemFree_bytes\n")
91
```

```
51 async def list_aggregates(stub: retrieval_pb2_grpc.RetrievalServiceStub,
52     request: retrieval_pb2.RetrievalRequest) -> None:
53     results = stub.ListAggregates(request)
54     async for result in results:
55         if result.name:
56             print(f"| nome metrica: {result.name} | partizione: {result.partition} | range: {result.range} | tipologia: {result.typology} | value: {round(result.value, 1)} | valore predetto: {round(result.predicted_value, 1)} |")
57     print("\n")
58
```

asynch_client.py cont. – esempio output:1

```
----- ListAggregates -----  
| nome metrica: memory_memFree_bytes | partizione: default | range: 1h | tipologia: max | value: 12012399616.0 | valore predetto: 0.0 |  
| nome metrica: memory_memFree_bytes | partizione: default | range: 1h | tipologia: min | value: 8310600192.0 | valore predetto: 0.0 |  
| nome metrica: memory_memFree_bytes | partizione: default | range: 1h | tipologia: mean | value: 10646500352.0 | valore predetto: 0.0 |  
| nome metrica: memory_memFree_bytes | partizione: default | range: 1h | tipologia: std_dev | value: 889110976.0 | valore predetto: 0.0 |  
| nome metrica: memory_memFree_bytes | partizione: default | range: 3h | tipologia: max | value: 11879600128.0 | valore predetto: 0.0 |  
| nome metrica: memory_memFree_bytes | partizione: default | range: 3h | tipologia: min | value: 8948230144.0 | valore predetto: 0.0 |  
| nome metrica: memory_memFree_bytes | partizione: default | range: 3h | tipologia: mean | value: 10516400128.0 | valore predetto: 0.0 |  
| nome metrica: memory_memFree_bytes | partizione: default | range: 3h | tipologia: std_dev | value: 894534976.0 | valore predetto: 0.0 |  
| nome metrica: memory_memFree_bytes | partizione: default | range: 12h | tipologia: max | value: 11392600064.0 | valore predetto: 0.0 |  
| nome metrica: memory_memFree_bytes | partizione: default | range: 12h | tipologia: min | value: 9631039488.0 | valore predetto: 0.0 |  
| nome metrica: memory_memFree_bytes | partizione: default | range: 12h | tipologia: mean | value: 10618000384.0 | valore predetto: 0.0 |  
| nome metrica: memory_memFree_bytes | partizione: default | range: 12h | tipologia: std_dev | value: 899774016.0 | valore predetto: 0.0 |  
| nome metrica: memory_memFree_bytes | partizione: default | range: 10m | tipologia: max | value: 0.0 | valore predetto: 10221400064.0 |  
| nome metrica: memory_memFree_bytes | partizione: default | range: 10m | tipologia: min | value: 0.0 | valore predetto: 10138699776.0 |  
| nome metrica: memory_memFree_bytes | partizione: default | range: 10m | tipologia: mean | value: 0.0 | valore predetto: 10164599808.0 |
```

Data Retrieval - async_server.py

```
102 async def serve() -> None:
103     server = grpc.aio.server()
104     retrieval_pb2_grpc.add_RetrievalServiceServicer_to_server(RetrievalService(), server)
105     listen_addr = '[:,]:50051'
106     server.add_insecure_port(listen_addr)
107     logging.info("Starting server on %s", listen_addr)
108     await server.start()
109     await server.wait_for_termination()
110
111
112 if __name__ == '__main__':
113     logging.basicConfig(level=logging.INFO)
114     asyncio.run(serve())
```

```
15 class RetrievalService(retrieval_pb2_grpc.RetrievalServiceServicer):
16     def __init__(self) -> None:
17         super().__init__()
18
19         # Connect to mysql
20         try:
21             self.cnx = sql.connect(
22                 user="user",
23                 password="password",
24                 host="db",
25                 database="test_dsbd")
26             self.cursor = self.cnx.cursor()
27         except sql.Error as e:
28             print(f"Error connecting to Mysql DB: {e}")
29             sys.exit(1)
30
31
32     def __del__(self) -> None:
33         self.cnx.close()
```

```
88 async def ListAggregates(
89     self, request:retrieval_pb2.RetrievalRequest,
90     context: grpc.aio.ServicerContext) -> AsyncIterable[retrieval_pb2.RetrievalAggregates]:
91     self.cursor.execute(request.query)
92     for result in self.cursor.fetchall():
93         yield retrieval_pb2.RetrievalAggregates(
94             name=result[1], partition=result[2], range=result[3],
95             typology=result[4], value=result[5], predicted_value=result[6])
96     yield retrieval_pb2.RetrievalAggregates(
97         name='', partition='', range='',
98         typology='', value=0, predicted_value=0)
```

SLA MANAGER REST API

GET -> <http://localhost:80/read.php>

POST -> <http://localhost:80/create.php>

POST -> <http://localhost:80/update.php>

POST -> <http://localhost:80/delete.php>

POST -> http://localhost:80/read_status.php

POST -> http://localhost:80/read_violation_1h.php

POST -> http://localhost:80/read_violation_3h.php

POST -> http://localhost:80/read_violation_12h.php

POST -> http://localhost:80/read_execution_time_12h.php

POST -> http://localhost:80/read_future_violation_10m.php

SLA MANAGER - read_violation_1h.php

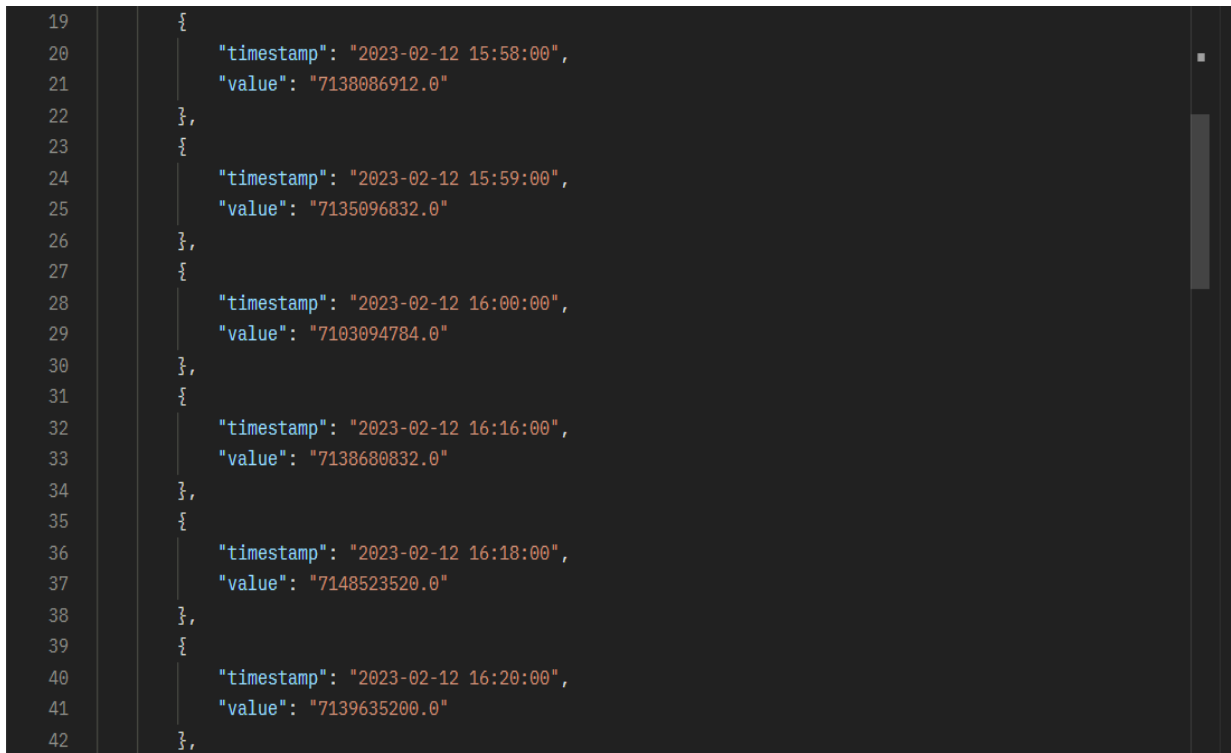
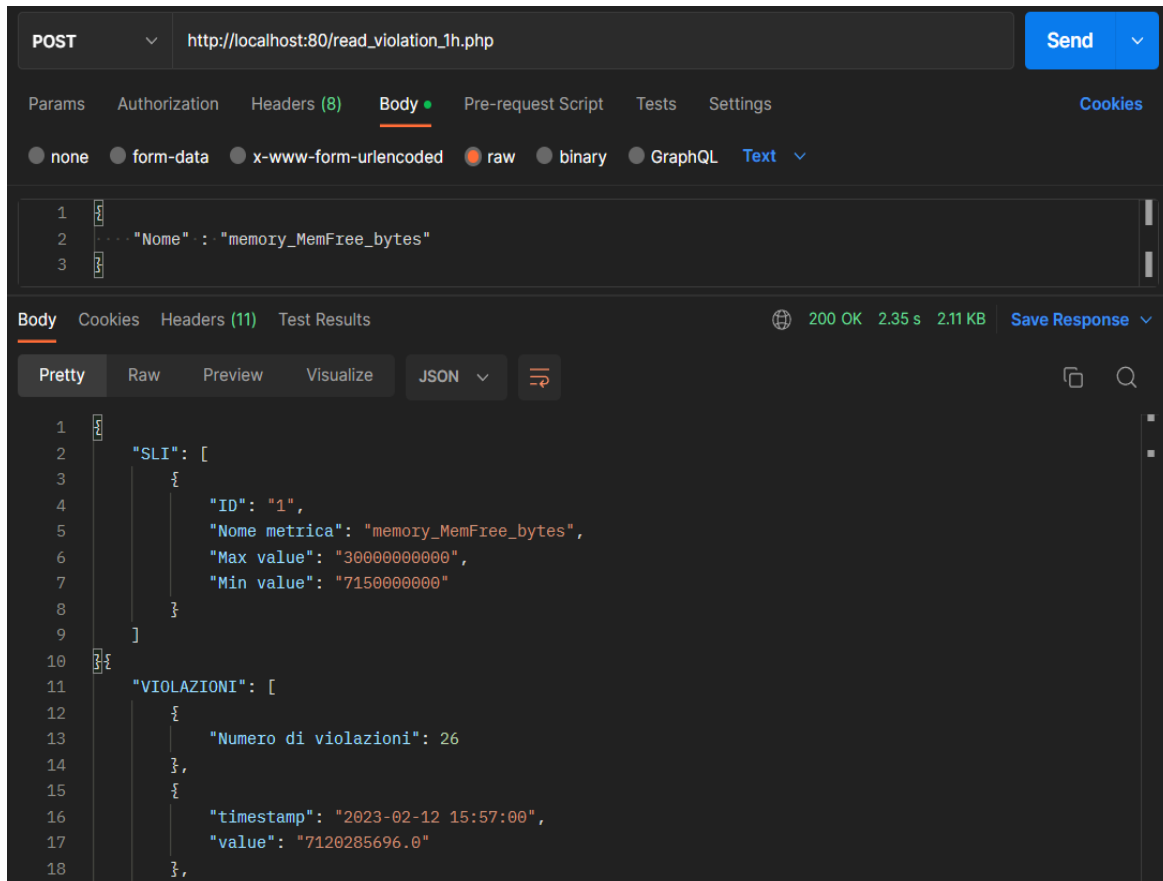
```
1 <?php
2 //headers
3 header("Access-Control-Allow-Origin: *");
4 header("Content-Type: application/json; charset=UTF-8");
5 header("Access-Control-Allow-Methods: POST");
6 header("Access-Control-Max-Age: 3600");
7 header("Access-Control-Allow-Headers: Content-Type, Access-Control-Allow-Headers, Authorization, X-Requested-With");
8
9 include_once '../config/database.php';
10 include_once '../models/SLA.php';
11
12 $database = new Database();
13 $db = $database->getConnection();
14 $sla = new SLA($db);
15
16 $data = json_decode(file_get_contents("php://input"));
17
18 $sla->Nome = $data->Nome;
19
20 // query products
21 $stmt = $sla->read_sli_status();
22 $num = $stmt->rowCount();
23
24 // se viene trovato l'SLI con quel nome nel database
25 if($num>0){
26     $arr = array();
27     $arr["SLI"] = array();
28     $row = $stmt->fetch(PDO::FETCH_ASSOC);
29     extract($row);
30     $item = array(
31         "ID" => $id,
32         "Nome metrica" => $nome,
33         "Max value" => $max_value,
34         "Min value" => $min_value,
35     );
36     array_push($arr["SLI"], $item);
```

```
37
38     $tempo = "1h";
39     $campioni = 60;
40     $command = 'python query_violation.py node_'. $sla->Nome. " ".$max_value. " ".$min_value. " ".$tempo. " ".$campioni;
41     $json = exec($command, $out, $status);
42     $array = json_decode($json, true);
43     $arr2 = array();
44     $arr2["VIOLAZIONI"] = array();
45     if(count($array)>0) {
46         $item2 = array(
47             "Numero di violazioni" => count($array)
48         );
49         array_push($arr2["VIOLAZIONI"], $item2);
50         foreach($array as $item)
51         {
52             $item3 = array(
53                 "timestamp" => $item['timestamp'],
54                 "value" => $item['value']
55             );
56             array_push($arr2["VIOLAZIONI"], $item3);
57         }
58         echo json_encode($arr);
59         echo json_encode($arr2);
60     }
61     else {
62         echo json_encode($arr);
63         echo json_encode(
64             array("message" => "Nessuna violazione")
65         );
66     }
67 }else{
68     echo json_encode($arr);
69     echo json_encode(
70         array("message" => "Nessun SLI Trovato avente tale nome")
71     );
72 }
73 ?>
```

SLA MANAGER - query_violation.py

```
1  import sys
2  from prometheus_api_client import PrometheusConnect, MetricRangeDataFrame
3  from prometheus_api_client.utils import parse_datetime
4  import json
5
6
7  client = PrometheusConnect(url='http://15.160.61.227:29090', disable_ssl=True)
8  label_config = {'job': 'host'}
9  start_time = parse_datetime(sys.argv[4])
10 #start_time = parse_datetime("1h")
11 end_time = parse_datetime("now")
12 metric_data = client.get_metric_range_data(
13     metric_name=sys.argv[1],
14     #metric_name="node_memory_MemAvailable_bytes",
15     label_config=label_config,
16     start_time=start_time,
17     end_time=end_time,
18 )
19 metric_df = MetricRangeDataFrame(metric_data)
20 if sys.argv[1] == 'node_filesystem_avail_bytes' or sys.argv[1] == 'node_filesystem_free_bytes':
21     data = metric_df.loc[metric_df['device']=='/dev/sda2', ['value']]
22 else:
23     data = metric_df.loc[:, ['value']]
24
25 sampled_metrics = data.resample('T').mean()
26 sampled_metrics = sampled_metrics.reset_index()
27 sampled_metrics = sampled_metrics.applymap(str)
28 sampled_metrics_dict = sampled_metrics.to_dict(orient='index')
29 max_value=float(sys.argv[2])
30 min_value=float(sys.argv[3])
31 interval=int(sys.argv[5])
32 #max_value=32000000000
33 #min_value=21840000000
34 for i in range (interval):
35     #for i in range (60):
36         current_value = float(sampled_metrics_dict[i]['value'])
37         if current_value < max_value and current_value > min_value:
38             del sampled_metrics_dict[i]
39
40 json_object = json.dumps(sampled_metrics_dict)
41 print(json_object)
```

SLA MANAGER - POSTMAN



Docker Compose

Ciascuno di questi servizi viene eseguito all'interno del proprio container. Per gestire la creazione e l'esecuzione di tutti i container è stato creato un unico file docker-compose.yml

```
version: '3'
services:
  prometheus:
    image: prom/prometheus
    ports:
      - "9090:9090"

  kafka:
    image: confluentinc/cp-kafka:latest
    depends_on:
      - zookeeper
    ports:
      - 9092:9092
    environment:
      KAFKA_BROKER_ID: 1
      KAFKA_ZOOKEEPER_CONNECT: zookeeper:2181
      KAFKA_ADVERTISED_LISTENERS: PLAINTEXT://kafka:9092,PLAINTEXT_HOST://localhost:29092
      KAFKA_LISTENER_SECURITY_PROTOCOL_MAP: PLAINTEXT:PLAINTEXT,PLAINTEXT_HOST:PLAINTEXT
      KAFKA_INTER_BROKER_LISTENER_NAME: PLAINTEXT
      KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR: 1
```

Docker Compose cont.

```
db:
  image: mysql:5.7
  environment:
    MYSQL_ROOT_PASSWORD: password
    MYSQL_DATABASE: test_dsbd
    MYSQL_USER: user
    MYSQL_PASSWORD: password
  volumes:
    - ./Database.sql:/docker-entrypoint-initdb.d/Database.sql
  restart: always

etl_data_pipeline:
  build:
    context: ./etl_data_pipeline
    dockerfile: Dockerfile
  depends_on:
    - prometheus
    - kafka
  environment:
    - PROMETHEUS_URL=http://15.160.61.227:29090
    - KAFKA_BOOTSTRAP_SERVERS=kafka:9092
    - KAFKA_TOPIC= 'prometheusdata'
```

```
data_storage:
  build:
    context: ./data_storage
    dockerfile: Dockerfile
  depends_on:
    - kafka
    - db
  environment:
    - KAFKA_BOOTSTRAP_SERVERS=kafka:9092
    - MYSQL_HOST=db
    - MYSQL_USER=user
    - MYSQL_PASSWORD=password
    - MYSQL_DB=test_dsbd

client_data_retrieval:
  build:
    context: ./data_retrieval/client
    dockerfile: Dockerfile
  stdin_open: true
  tty: true
  depends_on:
    - server_data_retrieval
```

```
server_data_retrieval:
  build:
    context: ./data_retrieval/server
    dockerfile: Dockerfile
  command: python3 /usr/app/async_server.py
  volumes:
    - ./data_retrieval/server:/usr/app/
  depends_on:
    - db
  environment:
    - MYSQL_HOST=db
    - MYSQL_USER=user
    - MYSQL_PASSWORD=password
    - MYSQL_DB=test_dsbd

sla_manager:
  build:
    context: ./sla_manager
    dockerfile: Dockerfile
  depends_on:
    - prometheus
    - db
  ports:
    - 80:80
  environment:
    - PROMETHEUS_URL=http://15.160.61.227:29090
    - MYSQL_HOST=db
```

Grazie per l'attenzione
